

Category, Question, Hints / Key Points, Difficulty PySpark, Explain RDD vs DataFrame vs Dataset, DataFrames are optimized, RDD is low-level, Dataset combines type safety and optimization, Medium PySpark, How to optimize PySpark jobs in Databricks?, Partitioning, caching, broadcast joins, avoiding shuffle, High PySpark, How do you handle skewed data in PySpark?, Salting keys, repartitioning, broadcast joins, High Python, Write a PySpark UDF for data transformation, UDF should be minimized due to performance, High Python, Handle missing/corrupt data in DataFrame, use dropna, fillna, or custom logic, Medium Delta Lake, Explain Delta Lake schema enforcement and evolution, Schema enforcement prevents bad data, evolution allows adding columns, Medium Delta Lake, How to perform incremental data loading in Databricks?, Use merge, upsert, or incremental filters, High ETL, Design a robust ETL pipeline in Databricks, Use modular notebooks, logging, error handling, High Scenario, Implement SCD Type 2 in PySpark, Use merge with conditions and version columns, High Scenario, Aggregate billions of rows efficiently in PySpark, Use proper partitioning, aggregations, avoid shuffles, Expert Advanced, Explain broadcast join vs shuffle join, Broadcast for small table, shuffle for large tables, High Advanced, Time travel in Delta Lake, Use versionAsOf or timestampAsOf to query older data, Medium Advanced, Handle schema changes in streaming pipelines, Use schema evolution and alter table, High Python, Combine multiple files from folder into DataFrame, Use spark.read.format(...).load(path), Medium Python, Profile and clean messy dataset, Use describe, dropna, fillna, regex, Medium PySpark, Window functions example, Rank, dense_rank, row_number with partitionBy, High ETL, Monitor job performance in Databricks, Use Spark UI, Ganglia metrics, explain plans, High Scenario, Find unmatched records between two datasets, Use join with indicator or anti-join, High Delta Lake, Optimize Delta tables using Z-ordering, Z-order for multi-column queries, High ETL, Handle late-arriving data or duplicates, Watermarks, deduplication, merge logic, High Advanced, Orchestrate notebooks in Databricks, Use Jobs API, workflows, triggers, High